

| Q.No | Question | Source |
|------|--|-------------------------------------|
| 1 | What happens during the Linux system boot process from power on to login prompt? | DEV.to [deadlock] |
| 2 | What steps occur when the <code>ls</code> command is issued in the terminal? | DEV.to [deadlock] |
| 3 | Explain what inodes are in Linux and their significance in the filesystem. | DEV.to [deadlock]; ZeroToMastery |
| 4 | What is the <code>/proc</code> filesystem and what information does it expose in Linux? | DEV.to [deadlock] |
| 5 | How would you troubleshoot a “filesystem full” error when <code>df</code> shows free space? | DEV.to [deadlock] |
| 6 | Name several performance monitoring tools you have used on Linux and describe their purpose. | DEV.to [deadlock] |
| 7 | What are the different types of filesystems supported on Linux and which have you used? | DEV.to [deadlock] |
| 8 | Explain kernel space vs user space and why this distinction is important. | DEV.to [deadlock] |
| 9 | What are processes and threads? How are they different? | DEV.to [deadlock]; ZeroToMastery |
| 10 | How does Linux manage process and thread scheduling? | ZeroToMastery |
| 11 | Explain Linux kernel memory management in brief. | DEV.to [deadlock] |
| 12 | What is the difference between stack and heap in process memory? | DEV.to [deadlock] |
| 13 | Describe race conditions and how Linux manages concurrency. | DEV.to [deadlock] |
| 14 | What does a high load average indicate? How do you investigate load average issues? | DEV.to [deadlock] |
| 15 | How would you troubleshoot high I/O issues on a Linux host? | DEV.to [deadlock] |
| 16 | What happens when the system runs out of inodes? | DEV.to [deadlock] |
| 17 | What is a zombie process? How do you find them and clean them up? | DEV.to [deadlock] |
| 18 | What is a runaway process and how would you identify it? | DEV.to [deadlock] |
| 19 | How do you change file permissions and ownership in Linux? Mention the commands. | ZeroToMastery |
| 20 | What is the purpose of the <code>chown</code> and <code>chmod</code> commands? | ZeroToMastery |
| 21 | What are sticky bits, setuid, and setgid? | ZeroToMastery |
| 22 | How do you manage and view running processes? Mention commands involved. | ZeroToMastery |
| 23 | How do you find and kill a process by name or port number in Linux? | ZeroToMastery |
| 24 | What is SSH? How do you secure SSH access for production servers? | LinkedIn: Praveen Singampalli |

| | | |
|----|--|-------------------------------------|
| 25 | How do you set limits for a user or a process in Linux? | ZeroToMastery |
| 26 | What is <code>nice</code> and <code>renice</code> ? How are they used? | ZeroToMastery |
| 27 | How do you delete or clean up empty files securely? | ZeroToMastery |
| 28 | Explain the differences and use-cases for hard links and symbolic links. | ZeroToMastery |
| 29 | How do you schedule recurring tasks on Linux? | ZeroToMastery |
| 30 | How would you monitor disk usage and find large files or directories quickly? | DEV.to [deadlock]; ZeroToMastery |
| 31 | What backup strategies and tools have you implemented on Linux? | DEV.to [deadlock] |
| 32 | How do you configure static and dynamic IP addresses on Linux interfaces? | ZeroToMastery |
| 33 | How would you set up and manage firewalls on a Linux server? | ZeroToMastery |
| 34 | What is SELinux/AppArmor and what role does it play in production infrastructure? | ZeroToMastery |
| 35 | How would you investigate and solve DNS resolution issues in Linux hosts? | DEV.to [deadlock] |
| 36 | What are the common troubleshooting steps when an application is unreachable over the network? | ZeroToMastery |
| 37 | How would you tune the Linux kernel parameters (<code>sysctl</code>)? | ZeroToMastery |
| 38 | What is shell scripting? Give an example of shell script automation in past work. | ZeroToMastery |
| 39 | How does package management differ between <code>apt</code> and <code>yum</code> ? | ZeroToMastery |
| 40 | Which log files do you frequently check for troubleshooting? Where are logs stored? | DEV.to [deadlock] |
| 41 | What is <code>systemd</code> ? How does it differ from <code>init</code> ? | DEV.to [deadlock] |
| 42 | Describe the sequence of actions when a user logs into a system via SSH. | DEV.to [deadlock] |
| 43 | How do you investigate and resolve “out of memory” issues in Linux? | DEV.to [deadlock] |
| 44 | What is <code>cgroups</code> ? How do you use <code>cgroups</code> to limit resource usage? | ZeroToMastery |
| 45 | Which commands or tools are useful for examining network usage and connections? | ZeroToMastery |
| 46 | What's your strategy for securing a freshly deployed public Linux server? | ZeroToMastery |
| 47 | What user and group management commands do you know? | ZeroToMastery |
| 48 | How do you perform kernel upgrades safely? | ZeroToMastery |
| 49 | How do you automate deployments or configuration management on Linux? | LinkedIn: Praveen Singampalli |

| | | |
|----|--|-------------------|
| 50 | Explain the basic workflow of troubleshooting a failed service on Linux. | DEV.to [deadlock] |
|----|--|-------------------|

Answers to the above questions.

| Q.No | Answer |
|------|---|
| 1 | During a Linux system boot: BIOS/UEFI initializes hardware, bootloader (GRUB/LILO) loads the kernel, which then initializes drivers. <code>init</code> or <code>systemd</code> starts essential services and user-space processes. Lastly, the login manager or shell prompts the user. |
| 2 | After <code>ls</code> is entered: The shell creates a process to run <code>/bin/ls</code> , which makes system calls to list directory contents, fetches metadata from inodes, and outputs to the terminal. |
| 3 | Inodes store filesystem metadata (owner, permissions, timestamps, disk block locations) for files/directories, but not names. |
| 4 | <code>/proc</code> is a virtual filesystem exposing kernel/process info (CPU, memory, process state) as files. |
| 5 | “Filesystem full” with space available: Check inode usage (<code>df -i</code>), disk quotas, files on hidden mount points, and open file descriptors holding deleted files (find via <code>lsof</code>). |
| 6 | Monitoring tools: <code>top</code> , <code>htop</code> , <code>vmstat</code> , <code>iostat</code> , <code>sar</code> , <code>dstat</code> , <code>free</code> , <code>nmon</code> are common, provide CPU, memory, IO, and process info. |
| 7 | Filesystems: <code>ext4</code> , <code>xfs</code> , <code>btrfs</code> , <code>zfs</code> , <code>nfs</code> , <code>tmpfs</code> , etc. <code>ext4/xfs</code> common for general use; <code>nfs</code> for network shares. |
| 8 | Kernel space: System and hardware code runs here, isolated from user space, where apps run. Separation prevents user bugs from crashing the system. |
| 9 | Process: Independent execution contexts with memory and resources. Thread: Lightweight, shares process memory/context. Threads within a process share most resources; processes are isolated. |
| 10 | Linux uses a scheduler (CFS in modern kernels) to allocate CPU time fairly among threads/processes based on priority and workload. |
| 11 | Linux manages memory with paging, virtual memory, cache, swap; using LRU and demand paging to optimize performance. |
| 12 | Stack: Stores function calls, local vars, grows/shrinks per call. Heap: Dynamically allocated at runtime with <code>malloc</code> , persists until freed. |
| 13 | Race conditions occur when processes/threads access shared resources concurrently. Use locks, semaphores, mutexes in Linux to avoid. |
| 14 | High load average indicates many runnable processes; check CPU, IO bottlenecks, run <code>top</code> , <code>uptime</code> , investigate blocking tasks. |
| 15 | Troubleshoot I/O: Use <code>iostat</code> , <code>iotop</code> , <code>vmstat</code> to check disk throughput, latency, IO wait states. Identify heavy writers/readers. |

| | |
|----|--|
| 16 | No inodes: Can't create new files, even with disk space. Check <code>df -i</code> , find/delete small files/directories. |
| 17 | Zombie process: Defunct, finished execution but not reaped by parent. Find with <code>` ps aux</code> |
| 18 | Runaway process: Consumes excess CPU/memory, often an infinite loop or bug. Identify with <code>top/ps</code> , <code>kill</code> or <code>debug</code> as needed. |
| 19 | Permissions: <code>chmod</code> (change mode), <code>chown</code> (change owner), <code>chgrp</code> (group). E.g., <code>chmod 755 file.txt</code> , <code>chown user:group file.txt</code> . |
| 20 | <code>chown</code> : Change file ownership (user/group). <code>chmod</code> : Change file/directory permissions (read/write/execute). |
| 21 | Sticky bit: Restrict file deletion to owner/root. <code>setuid/setgid</code> : Run file with file owner/group privileges (<code>chmod u+s/g+s file</code>). |
| 22 | View processes: <code>ps</code> , <code>top</code> , <code>htop</code> , <code>pgrep</code> , <code>pstree</code> , <code>jobs</code> (in shell). |
| 23 | Find/kill process by name: <code>pkill processname</code> ; by port: <code>fuser -n tcp PORTNUM</code> or <code>lsof -i :PORTNUM</code> , then <code>kill PID</code> . |
| 24 | SSH (Secure Shell) provides secure remote login. Secure by disabling root login, using key-based auth, changing default ports, restricting users, enabling fail2ban. |
| 25 | Set limits: User/per-process resource limits via <code>/etc/security/limits.conf</code> (<code>ulimit</code> for session), e.g. <code>nofile</code> , <code>nproc</code> , <code>memlock</code> . |
| 26 | <code>nice</code> : Set process priority at start. <code>renice</code> : Change priority of running processes. Lower value = higher priority. |
| 27 | Delete empty files: <code>find . -type f -empty -delete</code> . Secure delete: <code>shred</code> , <code>wipe</code> , or <code>srm</code> . |
| 28 | Hard link: Another name for same inode; symbolic (soft) link: pointer to filename. Hard links can't span filesystems or link directories. |
| 29 | Recurring tasks: Use <code>cron</code> (<code>crontab -e</code>), or <code>at</code> for one-time jobs. |
| 30 | Find large files: <code>du -h --max-depth=1</code> , <code>find . -size +100M -print</code> , <code>ncdu</code> , or <code>df -h</code> for space usage. |
| 31 | Backup strategies: <code>rsync</code> , <code>tar</code> , <code>cpio</code> , <code>dd</code> , automated <code>cron</code> , cloud backup solutions (e.g. AWS S3, <code>rclone</code>). Incremental and full backups. |
| 32 | IP configuration: Edit <code>/etc/network/interfaces</code> (Debian) or <code>ifcfg-*</code> (RedHat), use <code>ip addr</code> , <code>ifconfig</code> , <code>nmcli</code> or <code>netplan</code> for newer systems. |
| 33 | Firewalls: Use <code>iptables</code> , <code>firewalld</code> , <code>ufw</code> (Ubuntu-friendly). Define/allow/deny by ports/IPs. |
| 34 | SELinux/AppArmor: Offer mandatory access control, restricting program abilities for defense-in-depth in production Linux systems. |
| 35 | DNS troubleshooting: Use <code>nslookup</code> , <code>dig</code> , <code>host</code> for lookups; check <code>/etc/resolv.conf</code> ; verify network routes/firewalls. |
| 36 | Unreachable app: Check process status, logs, firewall, port availability (<code>netstat/ss</code>), DNS, and health checks. |

| | |
|----|---|
| 37 | Tune kernel: Edit <code>/etc/sysctl.conf</code> or use <code>sysctl -w key=value</code> ; e.g., for networking, memory, file descriptors. |
| 38 | Shell scripts automate tasks. Example: update packages, rotate logs, schedule cron backups. Use Bash, sh, etc. |
| 39 | apt (Debian/Ubuntu): Uses deb packages; yum (RHEL/CentOS): Uses rpm packages, different dependency management, configuration. |
| 40 | Logs: Main: <code>/var/log/</code> (syslog, messages, auth.log, dmesg, journalctl). Check relevant logs for service/debug info. |
| 41 | systemd: Modern PID 1 init system, service manager; replaces older <code>init</code> with dependency management, parallel startup, logging. |
| 42 | SSH login: Client connects, authenticates via key/password, server spawns login shell, initializes user environment. |
| 43 | Out of memory: System may invoke OOM killer, killing large processes. Monitor with <code>vmstat</code> , <code>top</code> , check <code>/var/log/messages</code> , add swap, optimize memory use. |
| 44 | cgroups: Kernel feature to limit/resource groups of processes (CPU, memory, IO). Used by Docker, Kubernetes, systemd. |
| 45 | Network usage: <code>netstat</code> , <code>ss</code> , <code>iftop</code> , <code>nethogs</code> , <code>iptraf</code> , <code>tcpdump</code> , <code>wireshark</code> . |
| 46 | Securing Linux: Update, create non-root users, configure SSH, enable firewall, disable unused services, enforce password policies, audit logs. |
| 47 | User/group commands: <code>useradd</code> , <code>usermod</code> , <code>passwd</code> , <code>groupadd</code> , <code>groups</code> , <code>deluser</code> , <code>delgroup</code> . |
| 48 | Kernel upgrades: Use package manager (<code>apt</code> , <code>yum</code>), reboot, keep previous kernel(s) for fallback, test first in staging. |
| 49 | Automation/config mgmt: Ansible, Chef, Puppet, SaltStack, shell/Python scripts; CI/CD pipelines (Jenkins, GitLab CI), Docker, Kubernetes. |
| 50 | Service troubleshooting: Check status (<code>systemctl status</code>), logs (<code>journalctl</code>), config files, restart, diagnose dependencies/resources. |

Sources:

- [DEV.to SRE/DevOps Questions](#)
- [ZeroToMastery Linux Interview Prep](#)
- [LinkedIn: Praveen Singampalli](#)